


Lomb-Scargle: single calls and survey batches


Float64 computation across all four roles.

CPU competitor: best tested GPU competitor: nifty-ls cuvarbase PyPI 0.2.5 cuvarbase v1.0

One lightcurve · warm API latency



32 distinct lightcurves · measured batch time / 32



One A40 / Xeon Gold 6342 host; CPU quota 7.65 cores. Identical host inputs and full host output spectra.
Warm medians of 5; whiskers = observed range. CPU threads/workers tuned over 1, 4, 8. Synthetic survey-size arrays; data loading and detrending excluded.
“Best tested” is restricted to validated completed candidates; see selected_timings.csv and validation.json. N/A means no eligible completed measurement.