



LS batch structure changes the comparison

CPU competitor: best tested GPU competitor: nifty-ls cuvarbase PyPI 0.2.5 cuvarbase v1.0

Float64 throughput



cuvarbase default precision



32 synthetic TESS-size lightcurves: 20k observations × 13.5k frequencies. A40 / Xeon Gold 6342, CPU quota 7.65 cores.

Shared-time inputs permit nifty-ls native batching. Full host outputs; medians of 5, observed ranges shown.

Default panel: CPU float64, cuvarbase float32, fastest validated GPU competitor precision. Inputs are identical across methods within each workload.